

# Scaleable Computing

## About Passwords and Hashes

CS7NS1/CS4400

Stephen Farrell

[stephen.farrell@cs.tcd.ie](mailto:stephen.farrell@cs.tcd.ie)

<https://github.com/sftcd/cs7ns1/>

Note: PRs for repo are welcome!

# Contents

- Background on passwords and their uses
- Password handling and a bit on password cracking
- Non-content: HOWTO do the practical assignments
- Exercise for you today: listen, learn, ask questions, and think about how scaling (up and down) affects the things described
- Note on references: If there's a reference here that can be followed, then that material is fair game for a supplemental exam (and worth a look regardless)
  - I'll try include URLs for such, lemme know if I messed that up

# Aside: Github Pull Requests

- Pull requests on github repo really are welcome
  - Won't contribute to marks directly, sorry
  - Might be a small bit of CV-filler though
- If you spot something wrong in the materials
- If you have something you'd like to add that can be public or useful to other students this or other years
- Who knows how to do that?
  - It's a good thing (TM) to know

# Passwords are horrible (1)

- Weak (guessable) passwords are inevitable – implicit in allowing “human memorable” values
  - 123456 almost always the most-used
- Re-use of (related) passwords is inevitable
  - Das, Anupam, et al. "The Tangled Web of Password Reuse." NDSS. Vol. 14. 2014.  
<https://www.cs.cmu.edu/~anupamd/paper/NDSS2014.pdf>
- Attacker often does not need all passwords, just a few good ones
  - e.g. belonging to an admin

# Passwords are horrible (2)

- Password entry user interfaces lend themselves to phishing
  - Distinguishing application “chrome” from content is v. Hard
- Forced “hard” password policies or password rotation force users to game the system
- Accessibility? What if you can’t use a keyboard?

# But... passwords are also great!

- Can be human memorable
  - Try that with your SSH private key!
- No need for h/w or special s/w
- People can change passwords (even if they don't)
- Can link between diverse systems
  - E.g. IPsec VPN access governed by ActiveDirectory login
- Fallback is required in any system – and you can fallback to closing a loop for a password reset
  - Main fail of a scheme I helped with called HOBA – RFC 7468
    - <https://tools.ietf.org/html/rfc7486>

# But... passwords leak (1)

- Attackers very rarely, if ever, try to grab passwords as they go by in a network packet (but they might, so don't allow that)
  - Careful of man-on-the-side attacks!
- Attack password entry (phishing, keylogging)
  - Can be done at scale, if you can get malware to host
- Masquerade as server, e.g. via borked DNS in coffee shop
- Grab a copy of the password verifier database
  - And then dictionary attack that – 1% success rate can be enough!

# But... passwords leak (2)

- Trickle brute-force attacks on online services
  - Esp. sshd but anything really – that happens **all** the time for all services exposed to the Internet
- Access/purchase a DB of leaked passwords and try those elsewhere or on the original leaky site
- Bugginess: user enters password in username field; system logs failed login by user “123456”; “oops!” says user, then successfully logs in as “joeblow”; system logs get centralised; logs contain cleartext passwords easily correlated with usernames
- How else might passwords leak?



# 2008 count of my passwords

Category	Count	Known	Examples
Logins	10	6	Laptops, host systems (incl. root accounts)
Devices	5	0	DSL router, home print/file servers, sensor nodes
Network access	4	0	Work n/w, WLAN, ISP, etc
Protocol	14	1	Outbound HTTP proxy, IMAP, Jabber, skype, etc.
Service	21	4	Mainly web sites with password stored outside browser
PINs	7	4	Bank cards, door access codes...
Total	61	15	

# Your passwords?

- How many? What kinds? How chosen? How protected? How long-lived?
- <your input here>

# What's a person to do? (1)

- Use password managers
  - Pros/cons?
- Avoid new accounts
  - Can you? Do you? Why? Why not?
- Avoid passwords where possible
  - SSH, USB tokens, even HOBA:-)
  - Key management?
- 2FA if you can and are willing/able

# What's a person to do? (2)

- Try hard to never use a terribly weak password
  - You never know if the system for which you create/update a password will end up becoming important for you
- Do not kick the bucket!
  - Post-mortem credential handling is kinda hard (but a real issue)

# What's a sysadmin to do? (1)

- Try not use passwords, esp. for admin purposes
  - Does a user **really** need to create an a/c? Is the minimal marketing benefit worth the risk of adding possibly toxic data to your DB?
- Try hard to insulate users from client-side attacks: CORS, XSS, etc.
- Monitor and protect your systems
  - fail2ban, denyhosts, other intrusion detection systems (IDSes)
- Try (but fail) to deploy universal single-sign-on (SSO)

# What's a sysadmin to do? (2)

- Avoid holding the password verifier DB, e.g. outsource to mega-scaler or service provider
  - But – centralisation of the web is a real problem, as is control
- Ensure best-practices for password verifier DB
  - More on that in a minute
- Maybe: use password authenticated key exchange (PAKE) schemes where that makes sense
  - Caveat: I'm a skeptic for many claimed uses of PAKEs – the problems with passwords are **not** really cryptographic ones

# What's a sysadmin to do? (3)

- Use two-factor authentication (2FA)... which does help
- Need caution in ensuring authentications are really out-of-band of one another
  - SMS messages – SS7 protocol attacks: Holtmanns, Silke, and Ian Oliver. "SMS and one-time-password interception in LTE networks." Communications (ICC), 2017 IEEE International Conference on. IEEE, 2017. <https://ieeexplore.ieee.org/document/7997246/>
- Google USB token claiming success against phishing
  - <https://krebsonsecurity.com/2018/07/google-security-keys-neutralized-employee-phishing/>

# What's a sysadmin not to do?

- Do not follow “traditional” password policies
  - Read this instead: Florêncio, Dinei, Cormac Herley, and Paul C. Van Oorschot. "An Administrator's Guide to Internet Password Research." LISA. Vol. 14. 2014.  
<https://www.usenix.org/system/files/conference/lisa14/lisa14-paper-florencio.pdf>
  - But main result is there's a huge gap between online guessable and offline dictionary attackable – do require systems/user-pwds to not be online guessable but no more, and do make sure offline dictionary attacks (and hence passwords) are hard-enough but don't care much more than that
- Do not enforce password “quality” requirements
  - Do encourage (not require) higher-entropy passwords e.g. via meters: Egelman, Serge, et al. "Does my password go up to eleven?: the impact of password meters on password selection." Proceedings of the SIGCHI Conference on Human Factors in Computing Systems. ACM, 2013.  
<https://www.guanotronic.com/~serge/papers/chi13b.pdf>
- NIST recanted on decades-old guidance recently
  - <https://www.passwordping.com/surprising-new-password-guidelines-nist/>



# Biometrics eh?

- Beware of biometrics!
  - You cannot easily change your fingerprint, retina, toe-smell!
  - And biometric verifiers can be fooled as much as anything (and haven't been tested near as much as some things)
- CCC 2017: 55 minute video of breaking biometrics
  - [https://www.youtube.com/watch?annotation\\_id=annotation\\_2684251971&feature=iv&src\\_vid=plY6k4gvQsY&v=VVxL9ymiyAU](https://www.youtube.com/watch?annotation_id=annotation_2684251971&feature=iv&src_vid=plY6k4gvQsY&v=VVxL9ymiyAU)
- Fiebig, Tobias, Jan Krissler, and Ronny Hänsch. "Security Impact of High Resolution Smartphone Cameras." WOOT. 2014.
  - <https://www.usenix.org/system/files/conference/woot14/woot14-febig.pdf>

# Password verifiers

- User enters username (u) and password (p)
- System transmits u & p to verification system
  - Rarely sends a processed password – could do but system problems often mean p is in clear at the application layer even if encrypted via TLS at the transport layer – what system problems might those be?
- System uses u to get password verifier (pv) value from database (call that PVDB)
- System checks if  $f(p,pv) == 'ok'$  for some function 'f'
  - You need to be able to replace 'f' without changing all passwords so really PVDB contains some form: of: u,f,pv

# Salts

- PVDB has many u,f,pv entries
- What if two users have the same password?
- Don't want u1,f,pv1 and u2,f,pv1 in the same PVDB as that'd help an attacker quite a bit
  - New attack enabled: if I can read `“/etc/passwd/”` then I could keep changing my password until my pv is the same as someone else's
- So we use a salt – a randomly chosen string that's set whenever the password changes
  - Now we're storing some form of u,f,s,pv
- Longer salts make dictionary attacks harder
- One possibility is to use a cryptographic hash function to generate pv

# PVDB Protection

- Is it useful to encrypt the entire PVDB?
- What else might we do to protect the PVDB?
  - Hint: consider the risk as being something like:  
 $\text{probability(PVDB leaks)} \times \text{cost of the leak}$
- There are some PAKEs that allow one to distribute the PVDB in cryptographically clever ways, but AFAIK those aren't in real-world use
  - There was IPR on that IIRC

# Aside: Cryptographic Hashes (1)

- Cryptographic hashes are one-way functions that take arbitrary input and produce a fixed-size output
  - Reversing hash should be cryptographically “hard” - say  $2^{128}$  units of work for some sensible unit
- Many uses in cryptography, usually we want hash functions to be highly efficient
  - Signature on LARGE file...
- Hash functions should be:
  - Collision resistant: finding  $x, y$  s.t.  $H(x) = H(y)$  is hard
  - Pre-image resistant: given  $x$ , s.t.  $x = H(y)$  finding  $y$  is hard
  - 2nd-pre-image resistant: given  $x$ , finding  $y$  s.t.  $H(x) = H(y)$  is hard

# Aside: Cryptographic Hashes (2)

- Finding collisions is **much** easier than pre-images due to birthday paradox
- If hash output length is  $n$ , then  $2^{(n/2)}$  work to find a collision, and  $2^n$  to find a pre-image for a perfect hash function
  - It can be important to understand that difference
  - $2^{128}$  is a **lot** less than  $2^{256}$ !
- Hash functions: MD5, SHA-1, **SHA-256**, SHA-512, SHA-3 family, ...

# Properties of 'f'?

- Do we want 'f' to be speedy?
- Do we want 'f' to use as little memory as possible?
- Do we want 'f' to be easy to parallelise?
  
- 'f' is often called a password hashing algorithm but is really a verifier, the password hashing alg produces pv given p and s (and sometimes u)
  - It's ok to not be terminologically pure:-)
- See <https://password-hashing.net/> for much more on a recent competition related to that kind of function
  - Worth another aside on cryptographic competitions, but maybe keep that for the new year (in CS7NS3/CS4407)

# MD5

- MD5 is an old hash function, probably older than most in this room:-)
  - See RFC 1321 (<https://tools.ietf.org/html/rfc1321>)
  - MD5 has 128 bit output, is very fast and has been broken for collisions (with  $\sim 2^{18}$  work)
- But let's say we base f on a simple use of the MD5 hash function...
  - That's really dumb! But is done all the time.
- In that case 'f' is something like:
  - `(pv==md5(salt||p)?"ok":"not-ok")`
    - Where || is catenation
- Unsalted versions are also used (OMG!)
  - `(pv=md5(p)?"ok":"not-ok")`
- Really easy to reverse via dictionary attack



# Dictionary Attack (1)

- Say you have  $\text{md5}(x)$ , what's  $x$ ?
  - While MD5 is broken for collisions, pre-images are still much harder so we won't try reverse the hash directly (yet!)
- But we can guess  $x$ , esp if  $x$  is human memorable
  - Human memorable  $\Rightarrow$  40 bits or less of entropy  
 $\Rightarrow 2^{40}$  search space  $\Rightarrow$  easy
- Algorithm:
  - Define the search space;  $\text{guess} = \text{first\_guess}()$ ;
  - while  $f(\text{guess}, s, \text{pv}) \neq \text{"ok"}$ )  $\text{guess} = \text{next\_guess}()$ ;
- Note:  $\text{next\_guess}()$  might do more than just take the next word from a list, e.g. if current  $\text{guess} = \text{"password"}$   $\text{next\_guess}()$  might return  $\text{"password01"}$  or  $\text{"passw0rd"}$

# Dictionary Attack (2)

- Can be easier still if:
  - We know about the possible/likely alphabet
  - We have a set of substrings that may be (part of) the password
  - We know something about the length of the password
- Dictionary attacks: The set of algorithms that base guesses on a dictionary of words, guessing the next variant and keep going 'till we've run out of search space, computational resources or we've won the game.
- If the “words” used are just random strings from an alphabet less than some length then we'd call that a **brute force** attack
- If 'f' uses a salt, then we need to explore the space of 's' via brute force, but the space of 'p' could still use a dictionary
  - So the salt slows down the dictionary attack, in proportion to the length of the salt

# Rainbow tables (1)

- Another form of 'f' is to use 'p' as an encryption key so that  $pv=e(p, \text{"fixed-string"})$ 
  - Or  $pv=e(p||s, \text{"fixed-string"})$  or  $pv=e(p, s||\text{"fixed string"})$  or similar
- If  $pv=e(p, \text{"fixed-string"})$  then we can use a dictionary attack as always, but we could also leverage a time-memory trade-off using rainbow tables.
  - Oechslin, Philippe. "Making a faster cryptanalytic time-memory trade-off." Annual International Cryptology Conference. Springer, Berlin, Heidelberg, 2003.  
<https://lasec.epfl.ch/pub/lasec/doc/Oech03.pdf>

# Rainbow Tables (2)

- Size can be dozens to hundreds of GB
  - <https://project-rainbowcrack.com/table.htm>
- Rainbow tables are sets of chains of passwords where we only store the first and last elements of each chain
  - Next element in chain derived via password hash function and a 'reduction' function
- Searching: iterate reduction/hashing of 'pv' until you find a value in some chain, at that point 'p' is (likely) somewhere in the chain, and you can begin from the chain start until you find pv somewhere in the chain.

# So MD5 is crap, how's SHA-512?

- Standard linux hashes of the “sha-512” (or “\$6”) variety:
- See ``man 3 crypt`` on your local linux
- Uses SHA-512 hash, iterated 5000 times
  - ‘s’ and ‘p’ are chars from [a-zA-Z0-9/]
  - ‘s’ is 16 chars, default random per ‘p’
- Output from ``mkpasswd -m sha-512 foo`` has format:  
\$6\$<salt>\$<hash> and could be :  
\$6\$m45hbNT1w/f\$gSX6x7nnEwkYTWskeNmz.j7XALhcOVcLL/  
c5oxMfrZy4bbYZsKa2la2yQGPfs0zgSQnPjCtW3mDkPgVqTFqHh.
- Details are gnarly and may be described by:
  - <https://www.akkadia.org/drepper/SHA-crypt.txt>
  - Seems like a less clean PBKDF2 (RFC8018, <https://tools.ietf.org/html/rfc8018>)

# So we had crap, then messy, ...

- Unsalted MD5 is terrible because it's so quick
- Sha-512 flavour crypt is slow but attacks can benefit from parallelism, e.g. multiple GPUs
- Argon2 won the password hashing competition and is explicitly designed to be a better 'f' that's memory intensive and time-memory trade-off resistant
  - Internet-draft specification: <https://tools.ietf.org/html/draft-irtf-cfrg-argon2>
- All very nice, but after PHC win some issues discovered, “fixed” by Argon2d (Argon2i won the PHC), and deployment is (so far) v. limited

# A good password recovery paper...

- Hranický, Radek, Martin Holkovič, and Petr Matoušek. "On Efficiency of Distributed Password Recovery." Journal of Digital Forensics, Security and Law 11.2 (2016): 5.  
<https://commons.erau.edu/cgi/viewcontent.cgi?article=1380&context=jdfsl>
- Describes challenges/opportunities in distributed computation for hash cracking with a nice description of implementation and measurement issues
  - Should be a fine model for part of an end-of-module report
  - And should provide plenty of hints as to HOWTO approach practical assignments
  - But, don't take this as me saying "do what they did" your problems and the assets you can bring to bear to solve problems will (likely:-) differ

# Conclusion

- Passwords are unavoidable and crap
- They can be handled more or less well at a systems level and cryptographically
- Dictionary attacks (of various forms) will likely have some percentage success
- There are scaling and distributed computing issues in password cracking
- You'll have fun with the practical assignments!



# Questions? Things to add?

- <your text here>